

Assignment 1: Implementing the OpenMP “inc” Directive

The **FLUSH** directive identifies synchronization points at which the implementation must provide a consistent view of memory. It must appear at the precise point in the code at which the synchronization is required. To avoid flushing all variables, we specify a list. Thread-visible variables are written back to memory at the point at which this directive appears. Modifications to thread-visible variables are visible to all threads after this point. Subsequent reads of thread-visible variables fetch the latest copy of the data. The flush operation provides a guarantee of consistency between a thread's temporary view and memory.

Running `grep -iR flush` in the clang directory gives a number files. Out of which, the following are of importance:

```
modified: bindings/python/clang/cindex.py
modified: include/clang-c/Index.h
modified: include/clang/AST/OpenMPClause.h
modified: include/clang/AST/RecursiveASTVisitor.h
modified: include/clang/AST/StmtOpenMP.h
modified: include/clang/Basic/OpenMPKinds.def
modified: include/clang/Basic/StmtNodes.td
modified: include/clang/Sema/Sema.h
modified: include/clang/Serialization/ASTBitCodes.h
modified: lib/AST/OpenMPClause.cpp
modified: lib/AST/StmtOpenMP.cpp
modified: lib/AST/StmtPrinter.cpp
modified: lib/AST/StmtProfile.cpp
modified: lib/Basic/OpenMPKinds.cpp
modified: lib/CodeGen/CGOpenMPRuntime.cpp
modified: lib/CodeGen/CGOpenMPRuntime.h
modified: lib/CodeGen/CGStmt.cpp
modified: lib/CodeGen/CGStmtOpenMP.cpp
modified: lib/CodeGen/CodeGenFunction.h
modified: lib/Parse/ParseOpenMP.cpp
modified: lib/Sema/SemaOpenMP.cpp
modified: lib/Sema/TreeTransform.h
modified: lib/Serialization/ASTReaderStmt.cpp
modified: lib/Serialization/ASTWriterStmt.cpp
modified: lib/StaticAnalyzer/Core/ExprEngine.cpp
modified: tools/libclang/CIndex.cpp
```

The implementation of INC is similar to FLUSH, barring a few files. Therefore first we clone the flush code replacing flush with inc and then include the logic for increment. Here is the **COMMON IMPLEMENTATION of flush and inc** -

- `Basic/OpenMPKinds.def` defines the list of supported OpenMP directives and clauses.
The flush directive is defined as `OPENMP_DIRECTIVE(flush)` and the flush clause is defined as `OPENMP_CLAUSE(flush, OMPFlushClause)` in `Basic/OpenMPKinds.def`. I made similar definitions for INC.
- The Parse folder does the Syntax Analysis or Parsing.

- `Parse/ParseOpenMP.cpp`: If `Dkind` is `OMPD_flush`, and lookahead token is left parenthesis, then consume token and push copy of the current token back to stream to properly parse pseudo-clause `OMPFlushClause`. I wrote a similar case `OMPD_inc` below it.
- The `Sema` folder does the Semantic Analysis - Checks for errors that are not found in syntax analysis – Incompatibilities and mismatches, especially type checking.
- `ActOnOpenMPVarListClause()` is located in `Sema/SemaOpenMP.cpp` and calls `ActOnOpenMPFlushClause()` which returns `OMPFlushClause::Create(..)`.
- `StmtNodes.td`: The line `def OMPFlushDirective : DStmt<OMPExecutableDirective>` adds a Statement Node. Clang reads this file and generates a `StmtNodes.inc` file, which is used to define different statement classes and read by several classes to define their node visitor function. I define an `OMPIncDirective` node, which extends the `OMPExecutableDirective` class - a basic class for representing single OpenMP executable directive.
- `Sema/SemaOpenMP.h`: The function `ActOnOpenMPFlushDirective` is called on well-formed `#pragma omp flush` after parsing of the associated statement. I call the `OMPFlushClause::Create(Context, StartLoc, LParenLoc, EndLoc, VarList)`; from `ActOnOpenMPFlushDirective` function. We add our case `OMPD_flush` where we call the `ActOnOpenMPAllocatedDirective` function. A similar `OMPD_inc` has to be added.
- AST has code related to abstract syntax tree generation. `AST/StmtOpenMP.cpp` has the `OMPFlushDirective *OMPFlushDirective::Create()` which creates the directive. It takes the AST context, start and end location of directive and list of clauses as parameters. The class also has a method to create an empty directive with the place for specifies number of clauses (`NumClauses`). Similarly include `OMPIncDirective *OMPIncDirective::Create()` in this file.
- `RecursiveASTVisitor.h` - This file defines the `RecursiveASTVisitor` interface, which recursively traverses the entire AST.
- `StmtPrinter.cpp` - This file implements the `Stmt::dumpPretty/Stmt::printPretty` methods, which pretty print the AST back out to C code. We add our definition of `VisitOMPIncDirective()` here looking at `VisitOMPFlushDirective()`. `VisitOMPFlushDirective()` function in `AST/StmtPrinter.cpp` visits the given flush node in the tree and prints it.
- `StmtProfile.cpp` - This file implements the `Stmt::Profile` method, which builds a unique bit representation that identifies a statement/expression.
- `ASTReaderStmt.cpp` - Implements Statements and Expression deserialization. This implements the `ASTReader::ReadStmt` method.
- `ASTWriterStmt.cpp` - Implements serialization for Statements and Expressions. In `ASTWriterStmt.cpp` we define the function `VisitOMPIncDirective` preferably after the definition of `VisitOMPExecutableDirective`.
- `TreeTransform.h`: We also need to define a tree transformation in `TreeTransform.h` - This file implements a semantic tree transformation that takes a given AST and rebuilds it, possibly transforming some nodes in the process. Using `StmtNode.td`, this class will already declare the `TransformOMPIncDirective` function.

- `CodeGen/OpenMPRuntime.cpp`: The ``emitFlush()`` function calls `EmitRuntimeCall()` of class `CodeGenFunction`. It builds runtime call ``void __kmpc_flush(ident_t *loc)``. `__kmpc` functions are defined in the openmp repository.
- `Serialization/ASTBitCodes.h`: We need to create a record for our statement in the `enum StmtCode`. These constants describe the records that describe statements or expressions. These records occur within type and declarations block, so they begin with record values of 128. Each constant describes a record for a specific statement or expression class in the AST. To add our own record we modify the `StmtCode` enum in `ASTBitCodes.h` file to add `STMT_OMP_INC_DIRECTIVE`.

Additional files to be changed for INC directive -

- `CodeGen/CGExprScalar.cpp`: This file has a function ``EmitScalarPrePostIncDec()`` which contains implementation of unary increment/decrement operations on variables. Taking this as reference, I change the below file.
- `CodeGen/CGOpenMPRuntime.cpp`: I have placed the actual logic for the inc directive in the ``emitInc()`` function. The first element of the input array(`ArrayRef<const Expr*> arr`) is the lvalue. We add 1 to this lvalue and store the result back to source location(`SourceLocation Loc`).
- The integer type checking can be done by `(type->isIntegerType())`, where `type` is an instance of `QualType` class.

test.II: (Comparing IRs)

IR for unary post increment operation on a variable (a++)	IR generated for incrementing with inc pragma	Comments
<pre>... define i32 @main() #0 { %1 = alloca i32, align 4 store i32 0, i32* %1, align 4 %2 = load i32, i32* %1, align 4 %3 = add nsw i32 %2, 1 store i32 %3, i32* %1, align 4 ret i32 0 }</pre>	<pre>... define i32 @main() #0 { %1 = alloca i32, align 4 store i32 0, i32* %1, align 4 %2 = load i32, i32* %1, align 4 %3 = add i32 %2, 1 store i32 %3, i32* %1, align 4 %4 = load i32, i32* %1, align 4 %5 = call i32 @__printf(i8* getelementptr inbounds ([7 x i8], [7 x i8]* @.str, i32 0, i32 0), i32 %4) ret i32 0 }</pre>	<pre>//allocate mem for variable //set var = initial_value 0 //load var to %2 //add 1 to %2 and store in %3 //store %3 back to src location //remaining code for printf etc. after pragma</pre>

The table shows that both IRs match and thus the pragma works!

OUTPUT

```
skarnik@dn002:/gpfs/projects/CSE504/skarnik
[skarnik@dn002 skarnik]$ cat test.c
#include <stdio.h>

int main()
{
    int a = 0;
#pragma omp inc(a)
    printf("a: %d\n", a);
}

[skarnik@dn002 skarnik]$ clang -fopenmp test.c
[skarnik@dn002 skarnik]$ ./a.out
a: 1
[skarnik@dn002 skarnik]$
```

AST FOR INC DIRECTIVE

```
[skarnik@dn016 skarnik]$ clang -cc1 -ast-dump test.c
test.c:1:10: fatal error: 'stdio.h' file not found
#include <stdio.h>
      ^~~~~~
TranslationUnitDecl 0x453d090 <<invalid sloc>> <invalid sloc>
| -TypeDecl 0x453d5e0 <<invalid sloc>> <invalid sloc> implicit __int128_t '__int128'
|   | -BuiltinType 0x453d300 '__int128'
| -TypeDecl 0x453d648 <<invalid sloc>> <invalid sloc> implicit __uint128_t 'unsigned __int128'
|   | -BuiltinType 0x453d320 'unsigned __int128'
| -TypeDecl 0x453d918 <<invalid sloc>> <invalid sloc> implicit __NSConstantString 'struct __NSConstantString_tag'
|   | -RecordType 0x453d720 'struct __NSConstantString_tag'
|   |   | -Record 0x453d698 '__NSConstantString_tag'
| -TypeDecl 0x453d9b0 <<invalid sloc>> <invalid sloc> implicit __builtin_ms_va_list 'char *'
|   | -PointerType 0x453d970 'char *'
|   |   | -BuiltinType 0x453d120 'char'
| -TypeDecl 0x453dc78 <<invalid sloc>> <invalid sloc> implicit __builtin_va_list 'struct __va_list_tag [1]'
|   | -ConstantArrayType 0x453dc20 'struct __va_list_tag [1]' 1
|   |   | -RecordType 0x453da90 'struct __va_list_tag'
|   |   |   | -Record 0x453da00 '__va_list_tag'
| -FunctionDecl 0x453dd20 <test.c:3:1, line:9:1> line:3:5 main 'int ()'
|   | -CompoundStmt 0x4591a18 <line:4:1, line:9:1>
|   |   | -DeclStmt 0x45916b8 <line:5:3, col:12>
|   |   |   | -VarDecl 0x4591638 <col:3, col:11> col:7 used a 'int' cinit
|   |   |   |   | -IntegerLiteral 0x4591698 <col:11> 'int' 0
|   |   |   | -CallExpr 0x4591960 <line:7:3, col:22> 'int'
|   |   |   |   | -ImplicitCastExpr 0x4591948 <col:3> 'int (*)(const char *, ...)' <FunctionToPointerDecay>
|   |   |   |   |   | -DeclRefExpr 0x4591860 <col:3> 'int (const char *, ...)' Function 0x4591708 'printf' 'int (const char *, ...)'
|   |   |   |   |   | -ImplicitCastExpr 0x45919b0 <col:10> 'const char *' <BitCast>
|   |   |   |   |   |   | -ImplicitCastExpr 0x4591998 <col:10> 'char *' <ArrayToPointerDecay>
|   |   |   |   |   |   | -StringLiteral 0x45918c8 <col:10> 'char [7]' lvalue "a: %d\n"
|   |   |   |   |   |   | -ImplicitCastExpr 0x45919c8 <col:21> 'int' <LValueToRValue>
|   |   |   |   |   |   | -DeclRefExpr 0x45918f8 <col:21> 'int' lvalue Var 0x4591638 'a' 'int'
|   |   |   |   | -ReturnStmt 0x4591a00 <line:8:1, col:8>
|   |   |   |   |   | -IntegerLiteral 0x45919e0 <col:8> 'int' 0
|   |   | -FunctionDecl 0x4591708 <line:7:3> col:3 implicit used printf 'int (const char *, ...)' extern
|   |   | -ParmVarDecl 0x45917a0 <<invalid sloc>> <invalid sloc> 'const char *'
|   |   | -FormatAttr 0x4591808 <col:3> Implicit printf 1 2
```

References:

https://www.ibm.com/support/knowledgecenter/SSGH2K_12.1.0/com.ibm.xlc121.aix.doc/compiler_ref/prag_omp_flush.html
www.openmp.org/wp-content/uploads/openmp-4.5.pdf
<http://freecompilercamp.org/>